

Differentiable Neural Computers

HYBRID COMPUTING USING A NEURAL NETWORK WITH
DYNAMIC EXTERNAL MEMORY (GRAVES ET AL. 2016)

Konstantinos Kogkalidis

May 28, 2018

Logic and Computation

Overview: Probabilistic Programming

Cross-domain

- Data Flow Programming
- Bayesian Reasoning / Machine Learning
- Fuzzy Logic
- Functional Programming

Overview: Probabilistic Programming

Cross-domain

- Data Flow Programming
- Bayesian Reasoning / Machine Learning
- Fuzzy Logic
- Functional Programming

Intuition: *"Rather than explicitly write a program, write some **constraints** on the behavior of the desired program and use computational tools to search the program space for **models** satisfying these constraints."*

Overview: Probabilistic Programming

Cross-domain

- Data Flow Programming
- Bayesian Reasoning / Machine Learning
- Fuzzy Logic
- Functional Programming

Intuition: *"Rather than explicitly write a program, write some **constraints** on the behavior of the desired program and use computational tools to search the program space for **models** satisfying these constraints."*

PROGRAM	MODEL
Discrete	Continuous
Deterministic	Stochastic
Static	Adaptive

Overview: DNC

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

Overview: DNC

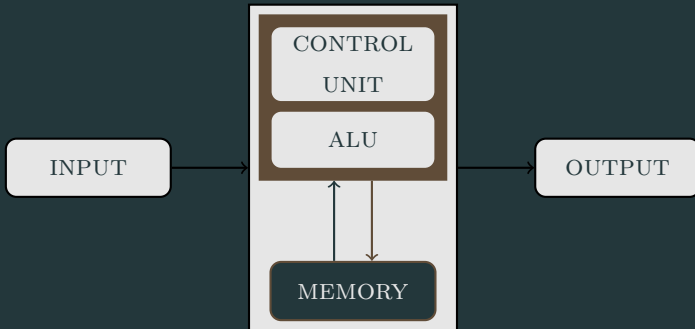
Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Functional replication of biological memory
- Reimaging of classical concepts of computation
- Extension of NTMs
 - End-to-end differentiable
 - Auto-associative memory
 - Turing complete
- + Computationally efficient memory management

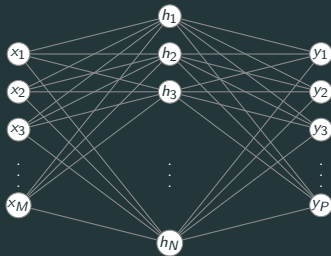
Introduction: Classic Computation

Von Neumann architecture



Introduction: Neural Networks

Simple Neural Net

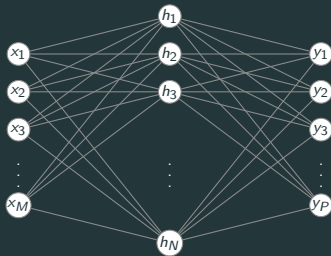


$$NN : \mathbf{x}_t \mapsto \mathbf{y}_t$$

No memory

Introduction: Neural Networks

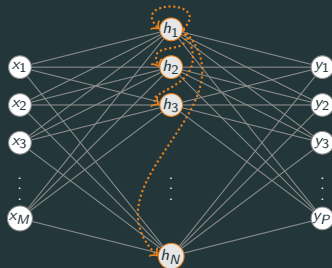
Simple Neural Net



$$NN : \mathbf{x}_t \mapsto \mathbf{y}_t$$

No memory

Simple Recurrent Net



$$RNN : \mathbf{x}_0 \otimes \mathbf{x}_1 \otimes \dots \otimes \mathbf{x}_t \mapsto \mathbf{y}_t$$

Finite memory

Introduction: Neural Networks

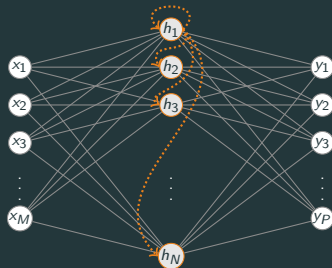
Simple Neural Net



$$NN : \mathbf{x}_t \mapsto \mathbf{y}_t$$

No memory

Simple Recurrent Net



$$RNN : \mathbf{x}_0 \otimes \mathbf{x}_1 \otimes \cdots \otimes \mathbf{x}_t \mapsto \mathbf{y}_t$$

Finite memory

*"If training vanilla neural nets is optimization over functions,
training recurrent nets is **optimization over programs**."*

Approach

Train a RNN to act as the **controller** of a memory matrix $M \in \mathbb{R}^{N \times W}$ through R **read heads** and one **write head**.

Approach

Train a RNN to act as the **controller** of a memory matrix $M \in \mathbb{R}^{N \times W}$ through R **read heads** and one **write head**.

1. Content Lookup

- **Attention** over memory defined by weightings $w \in \mathcal{S}^N$
- Compare controller output with memory objects (**auto-associative memory**)
- Allow partial matches (**pattern completion**)

Approach

Train a RNN to act as the **controller** of a memory matrix $M \in \mathbb{R}^{N \times W}$ through R **read heads** and one **write head**.

1. Content Lookup

- **Attention** over memory defined by weightings $w \in \mathcal{S}^N$
- Compare controller output with memory objects (**auto-associative memory**)
- Allow partial matches (**pattern completion**)

2. Sequential Retrieval

- Fill $L \in [0, 1]^{N \times N}$ indexing **temporal transitions**
- **Shift** operations defined by Lw , $L^\top w$

Approach

Train a RNN to act as the **controller** of a memory matrix $M \in \mathbb{R}^{N \times W}$ through R **read heads** and one **write head**.

1. Content Lookup

- **Attention** over memory defined by weightings $w \in \mathcal{S}^N$
- Compare controller output with memory objects (**auto-associative memory**)
- Allow partial matches (**pattern completion**)

2. Sequential Retrieval

- Fill $L \in [0, 1]^{N \times N}$ indexing **temporal transitions**
- **Shift** operations defined by Lw , $L^\top w$

3. Dynamic Allocation

- Mark memory locations with $u \in [0, 1]^N$ to **signal usage**
- Manipulate signals during R/W operations to enable **reallocation**
- Generalization to **unbounded memory**

Controller: Overview

A deep long short-term memory network receiving input:

$$\boldsymbol{\mathcal{X}}_t = [\boldsymbol{x}_t; \boldsymbol{r}_{t-1}^1; \dots; \boldsymbol{r}_{t-1}^R] \quad (\text{timestep } t)$$

and producing output:

$$(\boldsymbol{y}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \dots; \boldsymbol{\mathcal{X}}_t]; \boldsymbol{\vartheta}) \quad (\text{entire sequence})$$

where \mathcal{N} a set of state equations and $\boldsymbol{\vartheta}$ their trainable parameters.

LSTM: Overview

LSTM Network

Multiple stacked LSTM units.

LSTM Unit

A RNN that has an intrinsic memory cell c_t^I and three gates.

LSTM: Overview

LSTM Network

Multiple stacked LSTM units.

LSTM Unit

A RNN that has an intrinsic memory cell c_t^l and three gates.

1. Input gate i_t^l
2. Forget gate f_t^l
3. Output gate o_t^l

LSTM: Signal-Flow (1/2)

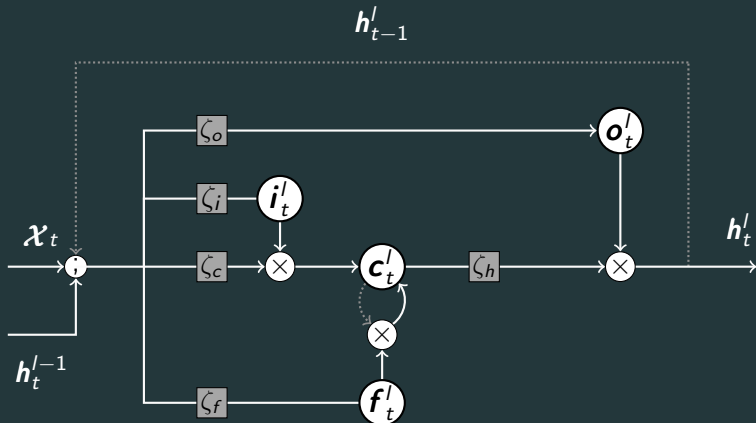
LSTM Unit (single layer)

- Input: $[\mathcal{X}_t; h'_{t-1}; h'^{-1}_t]$
- Output: h'_t

LSTM: Signal-Flow (1/2)

LSTM Unit (single layer)

- Input: $[\mathcal{X}_t; \mathbf{h}_{t-1}^l; \mathbf{h}_t^{l-1}]$
- Output: \mathbf{h}_t^l



Controller: Signal-Flow (2/2)

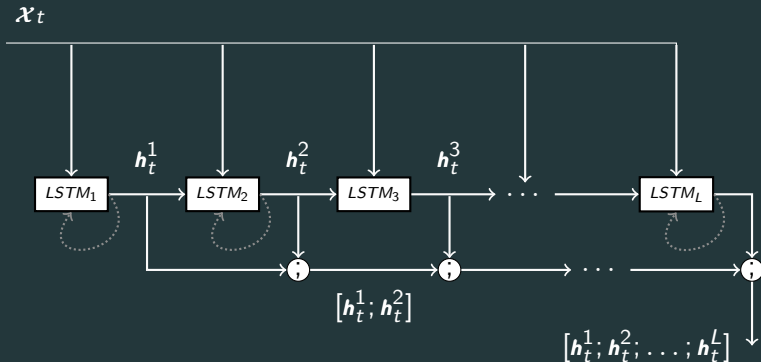
LSTM Network (multiple layers)

- Input: \mathcal{X}_t
- Output: $[h_t^1; h_t^2; \dots h_t^L]$

Controller: Signal-Flow (2/2)

LSTM Network (multiple layers)

- Input: \mathcal{X}_t
- Output: $[h_t^1; h_t^2; \dots; h_t^L]$



Controller: Outputs

$$(\mathbf{y}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\chi}_1; \dots; \boldsymbol{\chi}_T]; \boldsymbol{\vartheta})$$

Controller: Outputs

$$(\mathbf{y}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{x}_1; \dots; \boldsymbol{x}_T]; \boldsymbol{\vartheta})$$

User output $\mathbf{y}_t = W_y[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L]$

Controller: Outputs

$$(\mathbf{y}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{x}_1; \dots; \boldsymbol{x}_T]; \boldsymbol{\vartheta})$$

User output $\mathbf{y}_t = W_y[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L]$

Interface vector $\boldsymbol{\xi}_t = W_\xi[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L]$

- Read keys: $\mathbf{k}_t^{r,i}$
- Read strengths: $\beta_t^{r,i}$
- Write key: \mathbf{k}_t^w
- Write strength: β_t^w
- Erase vector: \mathbf{e}_t
- Write vector: \mathbf{v}_t
- Free gates: ϕ_t^i
- Allocation gate: g_t^a
- Write gate: g_t^w
- Read modes: π_t^i

Memory Addressing: Content-Lookup

R read keys $\mathbf{k}^{r,i} \in \mathbb{R}^W$, $i = 1 \dots R$

R read strengths $\beta^{r,i} \in [1, \infty)$, $i = 1 \dots R$

Write key $\mathbf{k}^w \in \mathbb{R}^W$

Write strength $\beta^w \in [1, \infty)$

Memory Addressing: Content-Lookup

R read keys $\mathbf{k}^{r,i} \in \mathbb{R}^W$, $i = 1 \dots R$

R read strengths $\beta^{r,i} \in [1, \infty)$, $i = 1 \dots R$

Write key $\mathbf{k}^w \in \mathbb{R}^W$

Write strength $\beta^w \in [1, \infty)$

Matching function \mathcal{D} comparing memory contents

Memory Addressing: Content-Lookup

R read keys $\mathbf{k}^{r,i} \in \mathbb{R}^W$, $i = 1 \dots R$

R read strengths $\beta^{r,i} \in [1, \infty)$, $i = 1 \dots R$

Write key $\mathbf{k}^w \in \mathbb{R}^W$

Write strength $\beta^w \in [1, \infty)$

Matching function \mathcal{D} comparing memory contents

Weighting function \mathcal{C} normalizing and sharpening matches

$$\mathcal{C}(M, \mathbf{k}, \beta)[i] = \frac{\exp\{\beta \mathcal{D}(\mathbf{k}, M[i, :])\}}{\sum_j \exp\{\beta \mathcal{D}(\mathbf{k}, M[j, :])\}}$$

Memory Addressing: R/W

Attention dictated by weightings $\mathbf{w} \in \mathcal{S}^N$

Erase vector $\mathbf{e}_t \in [0, 1]^W$

Write vector $\mathbf{v}_t \in \mathbb{R}^W$

Read operations

$$\mathbf{r}_t^i = M_t^\top \mathbf{w}_t^{r,i}$$

Memory Addressing: R/W

Attention dictated by weightings $\mathbf{w} \in \mathcal{S}^N$

Erase vector $\mathbf{e}_t \in [0, 1]^W$

Write vector $\mathbf{v}_t \in \mathbb{R}^W$

Read operations

$$r_t^i = M_t^\top \mathbf{w}_t^{r,i}$$

Write operations

$$M_t = \underbrace{M_{t-1} \circ (\mathbf{1} - \mathbf{w}_t^w \mathbf{e}_t^\top)}_{\text{erased memory}} + \underbrace{\mathbf{w}_t^w \mathbf{v}_t^\top}_{\text{new write}}$$

Memory Addressing: Dynamic Allocation

Free gates $\phi_t^i \in [0, 1]^W$

Allocation gate $g_t^a \in [0, 1]$

Write gate $g_t^w \in [0, 1]$

"Free list" scheme

$$\psi_t = \prod_{i=1}^R (1 - \phi_t^i w_{t-1}^{r,i}) \quad (\text{memory retention})$$

$$u_t = (\mathbf{u}_{t-1} + \mathbf{w}_{t-1}^w - \mathbf{u}_{t-1} \circ \mathbf{w}_{t-1}^w) \circ \psi_t \quad (\text{usage tracking})$$

Memory Addressing: Dynamic Allocation

Free gates $\phi_t^i \in [0, 1]^W$

Allocation gate $g_t^a \in [0, 1]$

Write gate $g_t^w \in [0, 1]$

"Free list" scheme

$$\psi_t = \prod_{i=1}^R (1 - \phi_t^i w_{t-1}^{r,i}) \quad (\text{memory retention})$$

$$u_t = (\mathbf{u}_{t-1} + \mathbf{w}_{t-1}^w - \mathbf{u}_{t-1} \circ \mathbf{w}_{t-1}^w) \circ \psi_t \quad (\text{usage tracking})$$

Attention shift

- Obtain the allocation vector \mathbf{a}_t by normalizing \mathbf{u}_t
- Shift \mathbf{w}_t by $g_t^a \mathbf{a}_t$ and scale by g_t^w

Memory Addressing: Temporal Linking

Read modes: $\pi_t^i \in \mathcal{S}_3$

Memory Addressing: Temporal Linking

Read modes: $\pi_t^i \in \mathcal{S}_3$

Temporal Transition $L_t \in [0, 1]^{N \times N}$

$$L[i, j] = \underbrace{(1 - w_t^W[i] - w_t^W[j])L_{t-1}[i, j]}_{\text{Part of last transition}} + \underbrace{w_t^W[i]w_{t-1}^W[j]}_{\text{Current transition}}$$

Memory Addressing: Temporal Linking

Read modes: $\pi_t^i \in \mathcal{S}_3$

Temporal Transition $L_t \in [0, 1]^{N \times N}$

$$L[i, j] = \underbrace{(1 - w_t^W[i] - w_t^W[j])L_{t-1}[i, j]}_{\text{Part of last transition}} + \underbrace{w_t^W[i]w_{t-1}^W[j]}_{\text{Current transition}}$$

Mode Interpolation

$$\mathbf{w}_t^{r,i} = \underbrace{\pi_t^i[1]L\mathbf{w}_t^{r,i}}_{\text{Forward shift}} + \underbrace{\pi_t^i[2]\mathbf{w}_t^{r,i}}_{\text{No shift}} + \underbrace{\pi_t^i[3]L^\top\mathbf{w}_t^{r,i}}_{\text{Backward shift}}$$

Experiments

- Language Tasks
 - Inference
 - Logical Reasoning

Experiments

- Language Tasks
 - Inference
 - Logical Reasoning
- Graph Tasks
 - Network Traversal
 - Policy Learning

Conclusion

Recap

- We can simulate computation and expand upon it by using differentiable, continuous operations.
- Simple programs (over complex data structures) can now be automagically inferred

Thank you!

Questions?

Appendix: LSTM Equations

$$i_t' = \sigma(W_i'[\mathbf{x}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_i') \quad (\text{input gate})$$

$$f_t' = \sigma(W_f'[\mathbf{x}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_f') \quad (\text{forget gate})$$

$$\mathbf{s}_t' = f_t' \mathbf{s}_{t-1}' + i_t' \tanh(W_s'[\mathbf{x}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_s') \quad (\text{state})$$

$$\mathbf{o}_t' = \sigma(W_o'[\mathbf{x}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_o') \quad (\text{output gate})$$

$$\mathbf{h}_t' = \mathbf{o}_t' \tanh(\mathbf{s}_t') \quad (\text{hidden})$$

$$\mathbf{v}_t = W_y[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L] \quad (\text{output vector})$$

$$\boldsymbol{\xi}_t = W_\xi[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L] \quad (\text{interface vector})$$

Appendix: Further Reading

Neural Architectures

- **Learning to Forget**
(Gers, Schmidhuber, Cummins)
- **Neural Turing Machines**
(Graves, Wayne, Danihelka)
- **Entity Networks**
(Henaff, Weston, Szlam, Bordes, LeCun)
- **End-to-End Memory Networks**
(Sukhbaatar, Szlam, Weston, Fergus)

- **Jointly Learning to Align and Translate**
(Bahdanau, Cho, Bengio)

Probabilistic Programming

- **Principles of Probabilistic Programming Languages**
(Goodman)
- **Backprop as a Functor**
(Fong, Spivak, Tuyras)
- **Formal Methods for Probabilistic Programming**
(Selsam, Liang, Dill)