

Differentiable Neural Computers

HYBRID COMPUTING USING A NEURAL NETWORK WITH
DYNAMIC EXTERNAL MEMORY (GRAVES ET AL. 2016)

Konstantinos Kogkalidis

May 28, 2018

Logic and Computation

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
 - End-to-end differentiable

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
 - End-to-end differentiable
 - Auto-associative memory

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
 - End-to-end differentiable
 - Auto-associative memory
 - Turing complete

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
 - End-to-end differentiable
 - Auto-associative memory
 - Turing complete
- + Memory attention mechanisms

Overview


Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
 - End-to-end differentiable
 - Auto-associative memory
 - Turing complete
 - + Memory attention mechanisms
- Mimic mammalian biological memory
- Employ classical concepts of computation


Introduction: Motivation

Von Neumann architecture



Introduction: Motivation

Simple Neural Net




$$y = g(h), \quad h = f(x)$$

No memory

Introduction: Motivation


Simple Neural Net



$$y = g(h), \quad h = f(x)$$

No memory

Recurrent Neural Net



$$h(t) = f([x(t); h(t-1)])$$

Finite, non-contiguous memory

Approach

Train a RNN to act as a **controller** to interact with a memory matrix of N (arbitrary many) addresses.

Approach

Train a RNN to act as a **controller** to interact with a memory matrix of N (arbitrary many) addresses.

1. Content Lookup

- **Attention** over memory defined by weightings $W \in \mathbb{R}^N$
- Compare controller output with memory objects
(**auto-associative memory**)
- Allow partial matches (**pattern completion**)

Approach

Train a RNN to act as a **controller** to interact with a memory matrix of N (arbitrary many) addresses.

1. Content Lookup

- **Attention** over memory defined by weightings $W \in \mathbb{R}^N$
- Compare controller output with memory objects
(**auto-associative memory**)
- Allow partial matches (**pattern completion**)

2. Sequential Retrieval

- Fill $L \in \{0, 1\}^{2N}$ indexing **temporal transitions**
- **Shift** operations defined by LW , $L^T W$

Approach

Train a RNN to act as a **controller** to interact with a memory matrix of N (arbitrary many) addresses.

1. Content Lookup

- **Attention** over memory defined by weightings $W \in \mathbb{R}^N$
- Compare controller output with memory objects
(**auto-associative memory**)
- Allow partial matches (**pattern completion**)

2. Sequential Retrieval

- Fill $L \in \{0, 1\}^{2N}$ indexing **temporal transitions**
- **Shift** operations defined by LW , $L^T W$

3. Dynamic Allocation

- Mark memory locations with $\{0, 1\}$ to **signal usage**
- Manipulate signals during R/W operations to enable **reallocation**
- Generalization to **unbounded memory**

Controller

A deep long short-term memory network receiving input:

$$\boldsymbol{\mathcal{X}}_t = [\mathbf{x}_t; \mathbf{r}_{t-1}^1; \dots; \mathbf{r}_{t-1}^R]$$

and producing output:

$$(\mathbf{v}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \dots; \boldsymbol{\mathcal{X}}_T]; \boldsymbol{\vartheta})$$

where \mathcal{N} a set of state equations and $\boldsymbol{\vartheta}$ their trainable parameters.

Controller: State Equations

A more detailed look into \mathcal{N} :

$$\mathbf{i}_t' = \sigma(W_i'[\boldsymbol{\chi}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_i') \quad (\text{input gate})$$

$$\mathbf{f}_t' = \sigma(W_f'[\boldsymbol{\chi}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_f') \quad (\text{forget gate})$$

$$\mathbf{s}_t' = \mathbf{f}_t' \mathbf{s}_{t-1}' + \mathbf{i}_t' \tanh(W_s'[\boldsymbol{\chi}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_s') \quad (\text{state})$$

$$\mathbf{o}_t' = \sigma(W_o'[\boldsymbol{\chi}_t; \mathbf{h}_{t-1}'; \mathbf{h}_t'^{-1}] + \mathbf{b}_o') \quad (\text{output gate})$$

$$\mathbf{h}_t' = \mathbf{o}_t' \tanh(\mathbf{s}_t') \quad (\text{hidden})$$

$$\mathbf{v}_t = W_y[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L] \quad (\text{output vector})$$

$$\boldsymbol{\xi}_t = W_\xi[\mathbf{h}_t^1; \dots; \mathbf{h}_t^L] \quad (\text{interface vector})$$

Controller: Signal-Flow

Single LSTM layer

