Differentiable Neural Computers

Hybrid Computing using a neural network with dynamic external memory (Graves et al. 2016)

Konstantinos Kogkalidis

May 28, 2018

Logic and Computation

Probabilistic Programming

Cross-domain

- Data Flow Programming
- Bayesian Reasoning
- Machine Learning
- Functional Programming

Probabilistic Programming

Cross-domain

- Data Flow Programming
- Bayesian Reasoning
- Machine Learning
- Functional Programming

Intuition: *"Rather than explicitly write a program, write some constraints on the behavior of the desired program and use computational tools to search the program space for models satisfying these constraints."*

Probabilistic Programming

Cross-domain

- Data Flow Programming
- Bayesian Reasoning
- Machine Learning
- Functional Programming

Intuition: *"Rather than explicitly write a program, write some constraints on the behavior of the desired program and use computational tools to search the program space for models satisfying these constraints."*

Program	Model
Discrete	*Continuous*
Deterministic	*Stochastic*
Static	*Adaptive*

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
    - End-to-end differentiable

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
  - End-to-end differentiable
  - Auto-associative memory

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
  - End-to-end differentiable
  - Auto-associative memory
  - Turing complete

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
  - End-to-end differentiable
  - Auto-associative memory
  - Turing complete
  + Memory attention mechanisms

Overview

Differentiable Neural Computer

A recurrent neural network coupled with an external memory.

- Extension of NTMs
  - End-to-end differentiable
  - Auto-associative memory
  - Turing complete
  + Memory attention mechanisms
- Mimic mammalian biological memory
- Employ classical concepts of computation

Introduction: Motivation

Von Neumann architecture

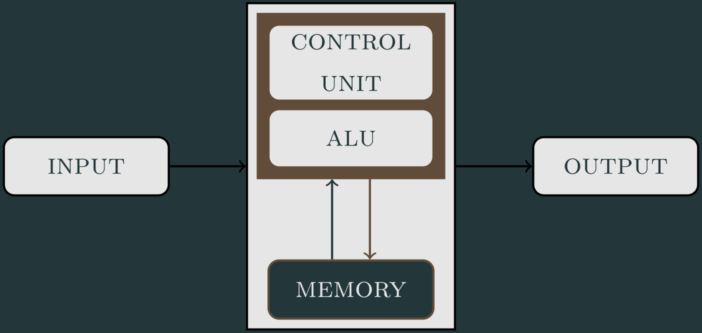

Introduction: Motivation

Simple Neural Net

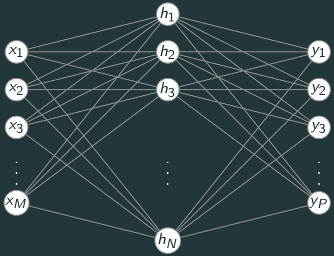

$$y = g(h),\ h = f(x)$$

No memory

Introduction: Motivation

Simple Neural Net

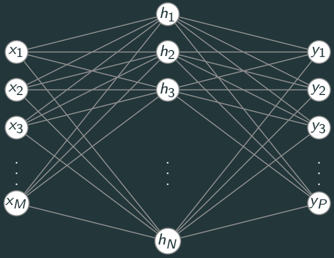

$y = g(h),\ h = f(x)$

No memory

Simple Recurrent Net

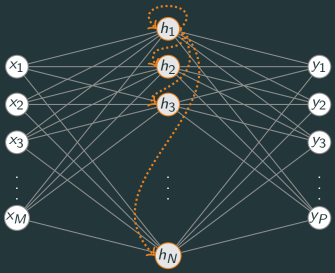

$h(t) = f([x(t); h(t-1)])$

Finite, non-contiguous memory

Introduction: Motivation

Simple Neural Net

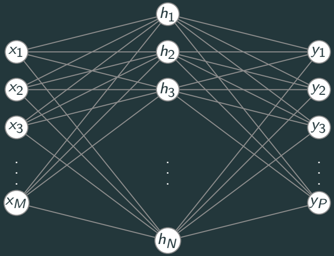

$$y = g(h),\ h = f(x)$$

No memory

Simple Recurrent Net

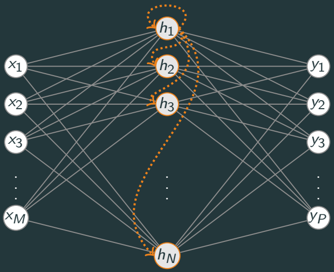

$$h(t) = f([x(t); h(t-1)])$$

Finite, non-contiguous memory

"If training vanilla neural nets is optimization over functions, training recurrent nets is optimization over programs."

Introduction: Motivation

Approach

Train a RNN to act as a controller to interact with a memory matrix of N (arbitrary many) addresses.

Approach

Train a RNN to act as a controller to interact with a memory matrix of N (arbitrary many) addresses.

1. Content Lookup
   - Attention over memory defined by weightings $W \in \mathbb{R}^N$
   - Compare controller output with memory objects (auto-associative memory)
   - Allow partial matches (pattern completion)

Approach

Train a RNN to act as a controller to interact with a memory matrix of N (arbitrary many) addresses.

1. Content Lookup
   - Attention over memory defined by weightings $W \in \mathbb{R}^N$
   - Compare controller output with memory objects (auto-associative memory)
   - Allow partial matches (pattern completion)
2. Sequential Retrieval
   - Fill $L \in \{0, 1\}^{2N}$ indexing temporal transitions
   - Shift operations defined by $LW, \ L^T W$

Approach

Train a RNN to act as a controller to interact with a memory matrix of N (arbitrary many) addresses.

1. Content Lookup
   - Attention over memory defined by weightings $W \in \mathbb{R}^N$
   - Compare controller output with memory objects (auto-associative memory)
   - Allow partial matches (pattern completion)
2. Sequential Retrieval
   - Fill $L \in \{0,1\}^{2N}$ indexing temporal transitions
   - Shift operations defined by $LW,\ L^T W$
3. Dynamic Allocation
   - Mark memory locations with $\{0,1\}$ to signal usage
   - Manipulate signals during R/W operations to enable reallocation
   - Generalization to unbounded memory

Controller

A deep long short-term memory network receiving input:

$$\boldsymbol{\mathcal{X}}_t = [\boldsymbol{x}_t; \boldsymbol{r}_{t-1}^1; \ldots; \boldsymbol{r}_{t-1}^R]$$

and producing output:

$$(\boldsymbol{v}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \ldots; \boldsymbol{\mathcal{X}}_T]; \vartheta)$$

where $\mathcal{N}$ a set of state equations and ϑ their trainable parameters.

Controller: State Equations

A more detailed look into $\mathcal{N}$:

$$
\begin{aligned}
\boldsymbol{i}_t^l &= \sigma(W_i^l[\boldsymbol{\mathcal{X}}_t; \boldsymbol{h}_{t-1}^l; \boldsymbol{h}_t^{l-1}] + \boldsymbol{b}_i^l) && \text{(input gate)} \\
\boldsymbol{f}_t^l &= \sigma(W_f^l[\boldsymbol{\mathcal{X}}_t; \boldsymbol{h}_{t-1}^l; \boldsymbol{h}_t^{l-1}] + \boldsymbol{b}_f^l) && \text{(forget gate)} \\
\boldsymbol{s}_t^l &= \boldsymbol{f}_t^l \boldsymbol{s}_{t-1}^l + \boldsymbol{i}_t^l tanh(W_s^l[\boldsymbol{\mathcal{X}}_t; \boldsymbol{h}_{t_1}^l; \boldsymbol{h}_t^{l-1}] + \boldsymbol{b}_s^l) && \text{(state)} \\
\boldsymbol{o}_t^l &= \sigma(W_o^l[\boldsymbol{\mathcal{X}}_t; \boldsymbol{h}_{t-1}^l; \boldsymbol{h}_t^{l-1}] + \boldsymbol{b}_o^l) && \text{(output gate)} \\
\boldsymbol{h}_t^l &= \boldsymbol{o}_t^l tanh(\boldsymbol{s}_t^l) && \text{(hidden)}
\end{aligned}
$$

$$
\begin{aligned}
\boldsymbol{v}_t &= W_y[\boldsymbol{h}_t^1; \ldots; \boldsymbol{h}_t^L] && \text{(output vector)} \\
\boldsymbol{\xi}_t &= W_\xi[\boldsymbol{h}_t^1; \ldots; \boldsymbol{h}_t^L] && \text{(interface vector)}
\end{aligned}
$$

Controller: Signal-Flow

Single LSTM layer

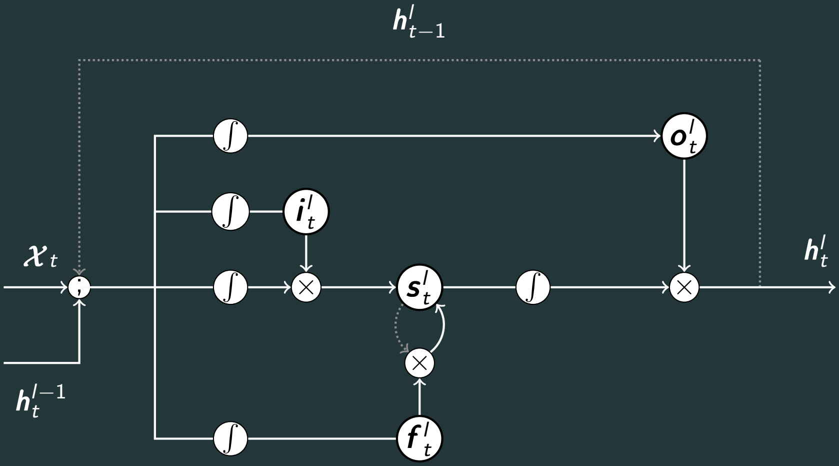

Controller: Outputs

$$(\boldsymbol{v}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \ldots; \boldsymbol{\mathcal{X}}_T]; \vartheta)$$

Controller: Outputs

$$(\boldsymbol{v}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \ldots; \boldsymbol{\mathcal{X}}_T]; \vartheta)$$

Output given by:

$$\boldsymbol{y}_t = \boldsymbol{v}_t + W_R[\boldsymbol{r}_t^1; \ldots; \boldsymbol{r}_t^R] \tag{hello}$$

Controller: Outputs

$$(\boldsymbol{v}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \ldots; \boldsymbol{\mathcal{X}}_T]; \vartheta)$$

Controller: Outputs

$$(\boldsymbol{v}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \ldots; \boldsymbol{\mathcal{X}}_T]; \vartheta)$$

$$\boldsymbol{\xi}_t = [\boldsymbol{k}_t^{r,1}; \ldots; \boldsymbol{k}_t^{r,R};\ \hat{\beta}_t^{r,1}; \ldots; \hat{\beta}_t^{r,R};\ \boldsymbol{k}_t^{w};\ \hat{\beta}_t^{w};$$
$$\hat{\boldsymbol{e}}_t;\ \boldsymbol{v}_t;\ \hat{\boldsymbol{f}}_t^{1}; \ldots; \hat{\boldsymbol{f}}_t^{R};\ \hat{g}_t^{a};\ \hat{g}_t^{w};\ \hat{\boldsymbol{\pi}}_t^{1}; \ldots; \hat{\boldsymbol{\pi}}_t^{R};]$$

Controller: Outputs

$$(\boldsymbol{v}_t, \boldsymbol{\xi}_t) = \mathcal{N}([\boldsymbol{\mathcal{X}}_1; \ldots; \boldsymbol{\mathcal{X}}_T]; \vartheta)$$

$$\boldsymbol{\xi}_t = [\overbrace{\boldsymbol{k}_t^{r,1}; \ldots; \boldsymbol{k}_t^{r,R};}^{\text{R read keys}}\ \hat{\beta}_t^{r,1}; \ldots; \hat{\beta}_t^{r,R};\ \boldsymbol{k}_t^{w};\ \hat{\beta}_t^{w};$$
$$\hat{\boldsymbol{e}}_t;\ \boldsymbol{v}_t;\ \hat{\boldsymbol{f}}_t^1; \ldots; \hat{\boldsymbol{f}}_t^R;\ \hat{g}_t^a;\ \hat{g}_t^w;\ \hat{\boldsymbol{\pi}}_t^1; \ldots; \hat{\boldsymbol{\pi}}_t^R;]$$